Mining Big Data in Real Time

Albert Bifet

YAHOO! RESEARCH

Turing/SLAIS 2012 Conference

BIG DATA

Measure and React

Motivation

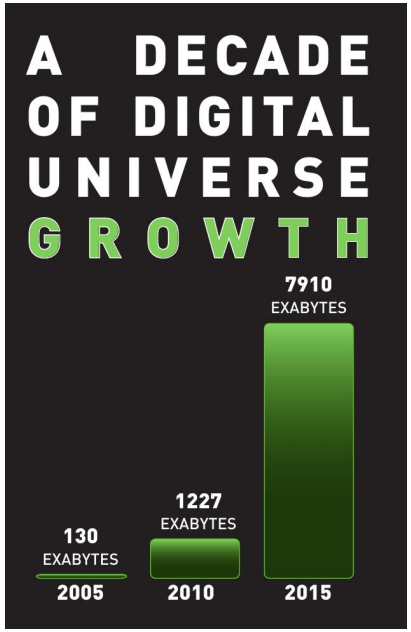

Source: IDC's Digital Universe Study (EMC), June 2011

Data is growing

Motivation

Memory unit	Size	Binary size
kilobyte (kB/KB)	10^3	2^{10}
megabyte (MB)	10^6	2^{20}
gigabyte (GB)	10^9	2^{30}
terabyte (TB)	10^{12}	2^{40}
petabyte (PB)	10^{15}	2^{50}
exabyte (EB)	10^{18}	2^{60}
zettabyte (ZB)	10^{21}	2^{70}
yottabyte (YB)	10^{24}	2^{80}

Data is growing

Motivation

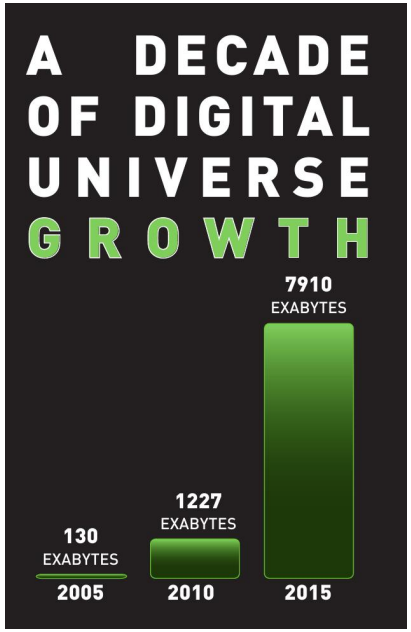

Source: IDC's Digital Universe Study (EMC), June 2011

Data is growing

Motivation

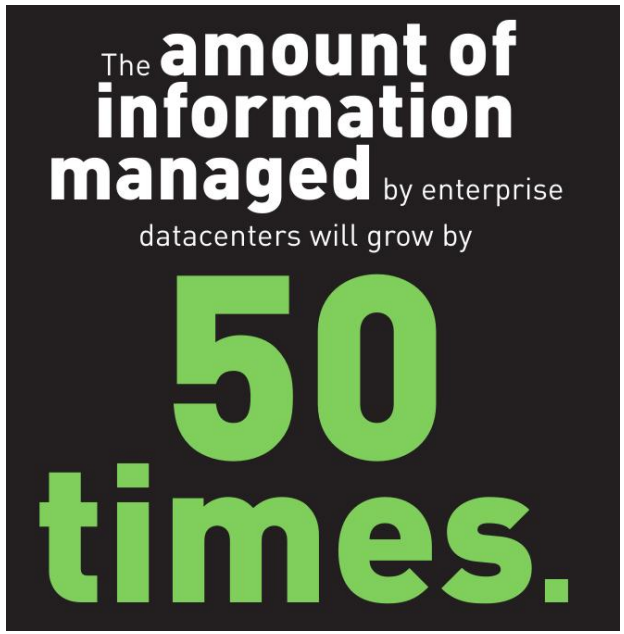

Source: IDC's Digital Universe Study (EMC), June 2011

Data is growing

Motivation

Source: IDC's Digital Universe Study (EMC), June 2011

Data is growing

Streaming Data

Big Data & Real Time

Big Data

McKinsey Global Institute (MGI) Report on Big Data, 2011.

Big data refers to datasets whose size is beyond the ability of typical database software tools to capture, store, manage, and analyze.

Big Data

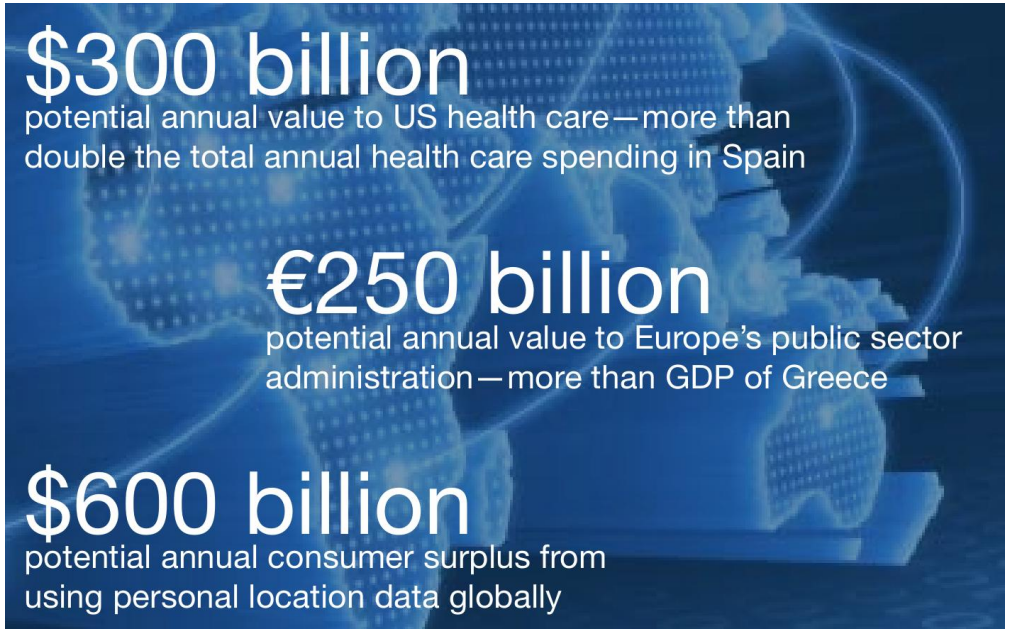

McKinsey Global Institute (MGI) Report on Big Data, 2011.

Big data refers to datasets whose size is beyond the ability of typical database software tools to capture, store, manage, and analyze.

BIG Data

- Volume
- Variety
- Velocity

3 Vs

Methodology

Sampling and distributed systems

Methodology

Paolo Boldi
Facebook Four degrees of separation

Big Data does not need big machines,
it needs big **intelligence**

Real time analytics

We want to analyze what is happening **now**.

Real time analytics

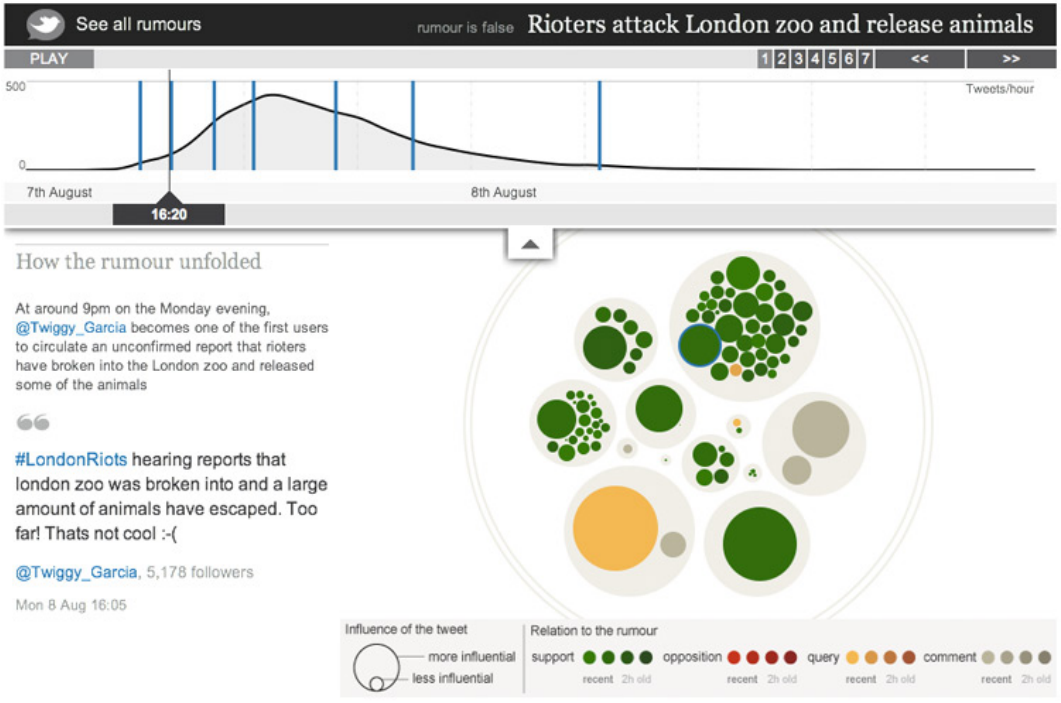

We want to analyze what is happening **now**.

Time and Memory

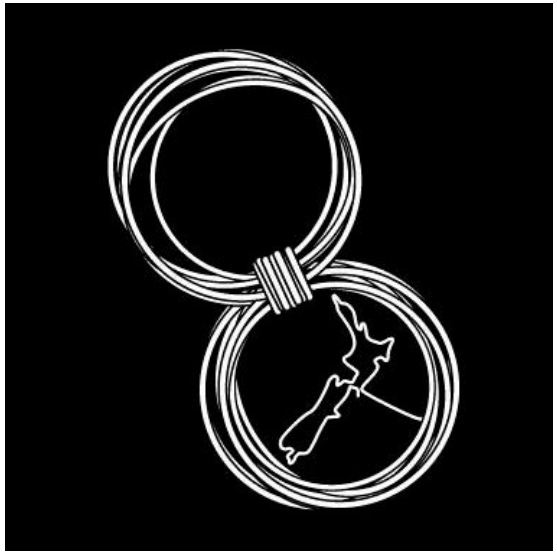

Number 8 Wire Mentality

Time and memory are the resource dimensions of the process.

Time and Memory

Time and memory are the resource dimensions of the process.

Algorithms

Classification, Regression, Clustering, Frequent Pattern Mining.

Applications

- sensor data: industry, cities
- telecomm data
- social networks: twitter, facebook, yahoo
- marketing: sales business

Data may come from: humans, sensors, or machines.

New applications: social networks

Twitter: A Massive Data Stream

- Micro-blogging service
- Built to discover what is happening at any moment in time, anywhere in the world.
- 3 billion requests a day via its API.

`MOA-TweetReader`: a real-time system to

- read tweets in real time
- detect changes
- find the terms whose frequency changed

Sentiment Analysis on Twitter

Sentiment analysis

Classifying messages into two categories depending on whether they convey positive or negative feelings

Emoticons are visual cues associated with emotional states, which can be used to define class labels for sentiment classification

Positive Emoticons	Negative Emoticons
:)	:(
:-)	:-(
:)	: (
:D	
=)	

Table : List of positive and negative emoticons.

New problem: structured classification

New methods for structured classification

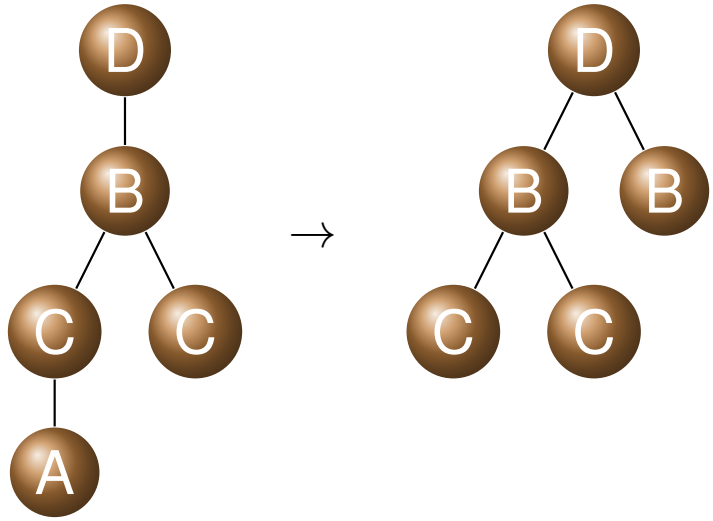

- sequences, trees, graphs

New problem: structured classification

New methods for structured classification

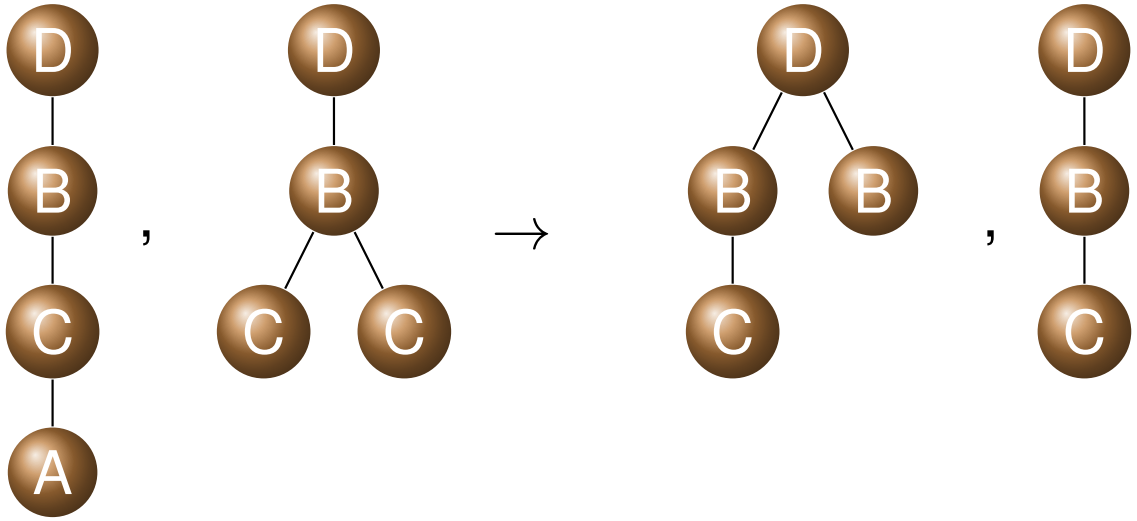

- sequences, trees, graphs
- frequent pattern mining techniques

New problem: structured classification

New methods for structured classification

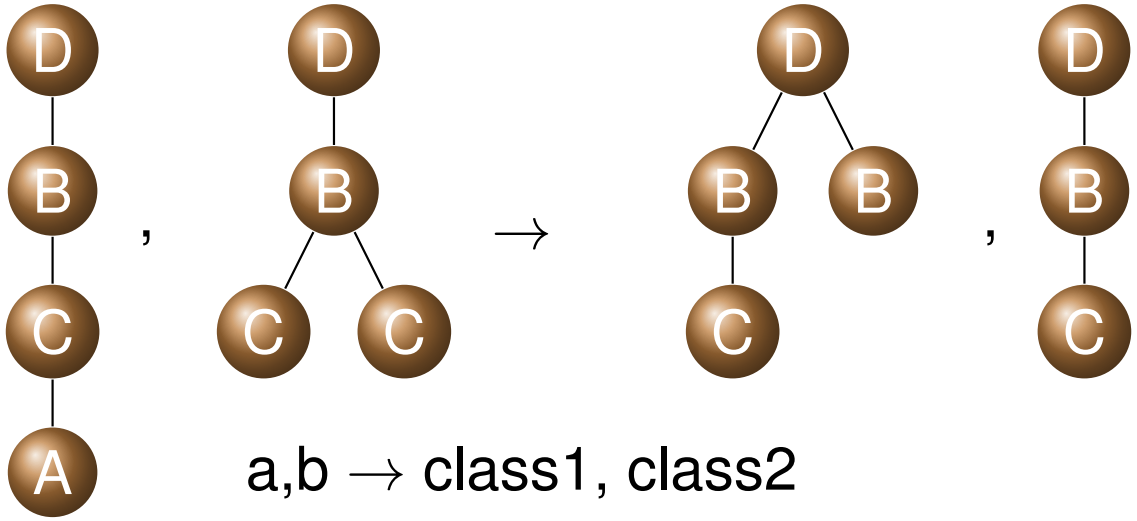

- sequences, trees, graphs
- frequent pattern mining techniques
- multi-label data mining
  - Example: Lord of the Rings → Action, Adventure, Fantasy

New Techniques: Distributed Systems

Hadoop, S4 and Storm

Hadoop

Hadoop

Hadoop

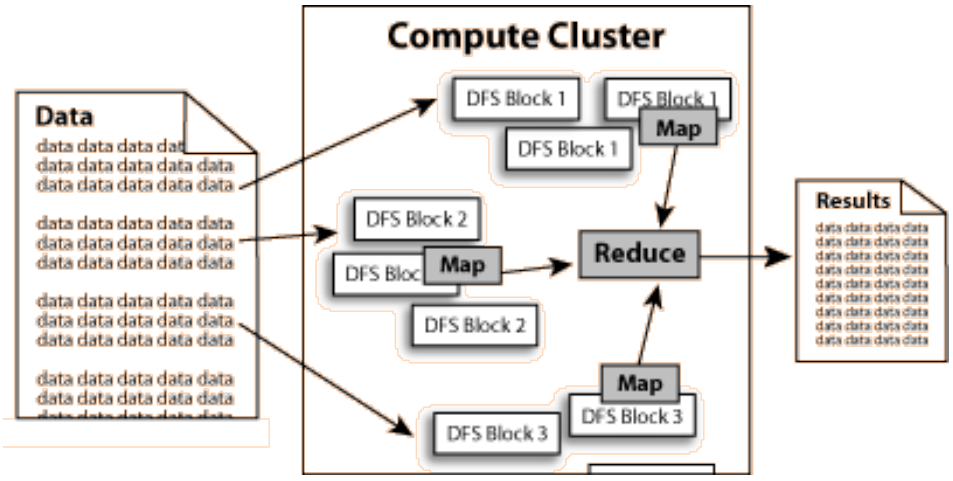

Hadoop architecture

Apache Mahout

Mahout: open source framework

Pig

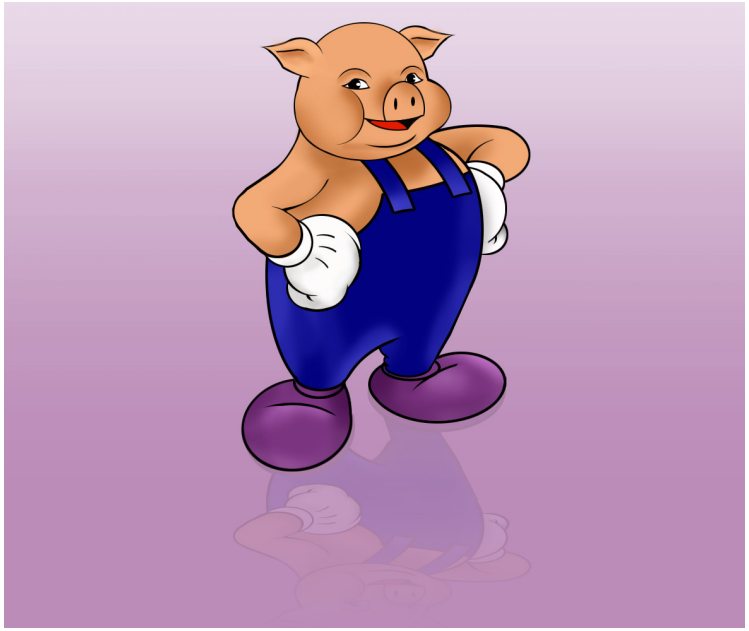

Pig: Similar to SQL

Pig

- `A = LOAD 'data' USING PigStorage() AS (f1:int, f2:int, f3:int);`
- `B = GROUP A BY f1;`
- `C = FOREACH B GENERATE COUNT ($0);`
- `DUMP C;`

Pig: Similar to SQL

Apache S4

Apache S4

Apache S4

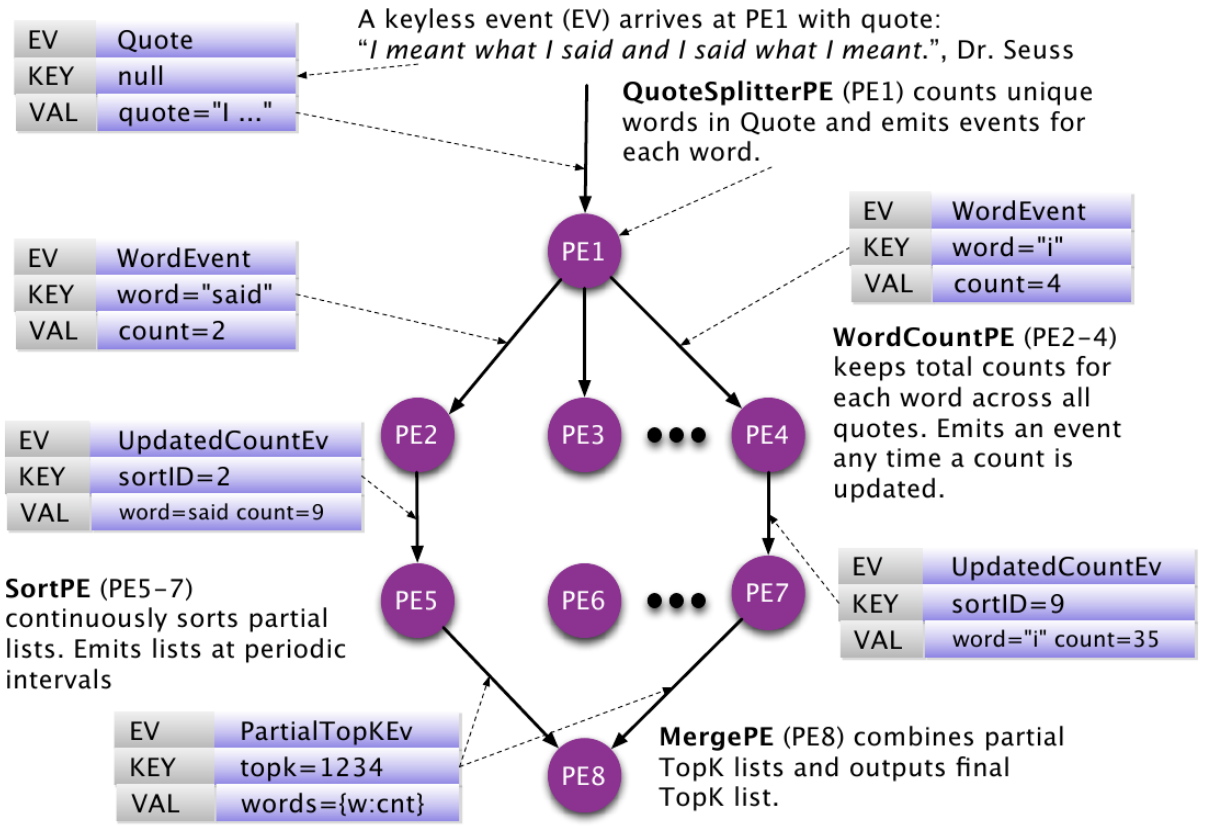

PE ID	PE Name	Key Tuple
PE1	QuoteSplitterPE	null
PE2	WordCountPE	word="said"
PE4	WordCountPE	word="i"
PE5	SortPE	sortID=2
PE7	SortPE	sortID=9
PE8	MergePE	topK=1234

Storm

Storm from Twitter

Storm

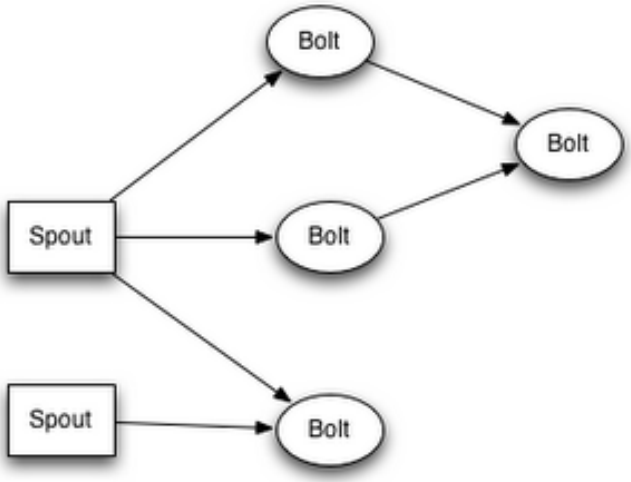

Stream, Spout, Bolt, Topology

Storm

Tools

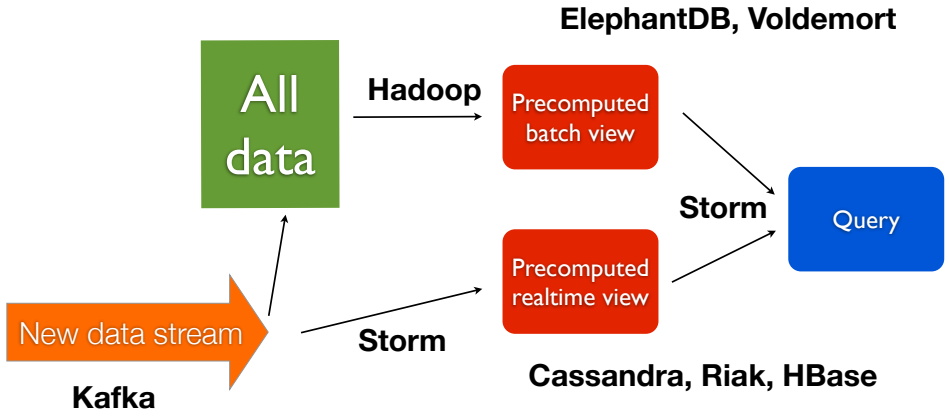

"Lambda Architecture"

Runaway complexity in Big Data
Nathan Marz, 2012

Data Streams

Big Data & Real Time

Data Streams

Thanks!